A Latent Oscillator Measurement Model for Simulating Emotion Dynamics in Video

Aleksandar Tomašević Hudson Golino Alexander P. Christensen

Abstract

Automated facial-expression scoring can turn video into dense time series of visible affect, but these channels are model outputs rather than direct observations of latent emotion dynamics. Behavioral applications therefore need simulation frameworks that make the latent measurement structure known before dynamic methods are interpreted substantively. We introduce the Latent Oscillator Measurement Model (LOMM), a generative framework for evaluating dynamic dimensional recovery in facial-expression-recognition-like video score series. LOMM generates second-order latent oscillator trajectories, gates their activity over time, and maps them through a block measurement layer into noisy observed score panels with known dimensionality, loading strength, measurement error, feedback regime, and activity pattern. We calibrated the generator against MAFW-derived score panels using targets for amplitude dispersion, local movement, temporal persistence, directional-regime complexity, spectral complexity, and dimensional spread. The calibrated generator reproduced several local temporal and dimensional signatures, with a remaining amplitude-dispersion mismatch. We then compared Dynamic Exploratory Graph Analysis (DynEGA), GraphicalVAR, and bounded GIMME in known-truth simulations, including community-detection choice as a diagnostic factor in the bounded comparator. DynEGA showed the strongest dimensional recovery and completed the expanded grid. GraphicalVAR recovered dimensionality in many conditions but was slower and more sensitive to weak measurement conditions, consistent with partial-correlation attenuation of common-cause measurement blocks. GIMME was more flexible but computationally less feasible in this design. LOMM links generative modeling, empirical realism checks, and known-truth recovery analysis so researchers can evaluate whether dynamic methods recover latent organization, tolerate measurement error, and scale to video-derived behavioral time series before substantive interpretation.

Introduction

Recent advancements in deep learning and computer vision have made it increasingly feasible to transform video frames into time series of emotional expression scores (S. Li & Deng, 2022; Meng et al., 2017). In a facial-expression recognition task, an algorithm maps each image frame or video segment to emotion-label predictions, usually as category probabilities or continuous scores for labels such as happiness, anger, sadness, fear, surprise, disgust, or neutral expression. For behavioral-science applications, these tools create a new kind of intensive longitudinal record: high-frequency, model-derived measurements of visible affective expression. These records differ substantially from self-report time series collected through ecological momentary assessment (EMA; Tong et al., 2005) or experience sampling methods (ESM; Hoemann et al., 2020). Facial-expression recognition (FER) scores are dense outputs of measurement systems applied to video frames, making them promising for studying the temporal organization of emotional expression at scale. Their use also raises a methodological problem: dynamic models need to separate latent expressive processes from the observed score series used to measure them.

The substantive motivation comes from video-rich behavioral settings in which visible affect is part of the behavior being studied. Facial expressions are central to human communication (Barrett et al., 2019; Mehu & Scherer, 2015; Stewart et al., 2009), but many behavioral-science questions concern how these expressions unfold over time rather than how a single frame should be labeled. In psychotherapy tasks, autism assessments, collaborative work, public communication, and social-media videos, the relevant questions are dynamic: whether expressions return quickly after perturbation, whether affective displays settle into stable or volatile regimes, and whether many model-output score channels reflect fewer latent expressive dimensions (Hollenstein, 2007; Houben et al., 2015; Kuppens & Verduyn, 2017).

Automated facial-analysis systems now make this kind of work feasible at scale. Datasets such as FER2013 (Goodfellow et al., 2013) supported deep learning approaches to facial-expression recognition, especially Convolutional Neural Networks (Arriaga et al., 2017; CNNs; Ezerceli & Eskil, 2022). Most early systems inherited their emotion categories from the training data and focused on fixed classification tasks, commonly returning probabilities for seven categories: anger, disgust, fear, happiness, sadness, surprise, and neutral facial expression (Wang et al., 2022). More recently, transformer-based and CLIP-like models have enabled zero-shot classification with user-specified emotion labels (Foteinopoulou & Patras, 2023; H. Li et al., 2023; Tomašević et al., 2026), making it possible to represent each video frame as a vector of emotion-label scores. Applied studies are already using such automated scoring in clinical, collaborative, and social-media settings (Halamová et al., 2023; Holiday et al., 2023; Q. Li et al., 2024; Manelis-Baram et al., 2025).

This technical progress also creates a validation problem. FER-like score series are model-derived, multivariate, and error-prone measurements of visible behavior. Their interpretation depends on whether the observed score channels are treated as emotion variables in their own right, as indicators of social interaction, or as noisy indicators of lower-dimensional latent affective processes (McKeown & Allison, 2024). For dynamic modeling, this distinction has practical consequences: measurement error can change the structure recovered by dynamic-network methods, making it important to separate latent temporal organization from the fallible score series used to observe it (Kankaanpää et al., 2026).

The aim of this paper is to introduce the Latent Oscillator Measurement Model (LOMM), a generative simulation model for FER-like emotion score series (Boker, 2015; Boker & Nesselroade, 2002; Ollero et al., 2023). LOMM first generates latent expressive trajectories from second-order oscillator dynamics and then maps those latent processes into FER-like classifier score outputs through an explicit measurement layer. In this sense, LOMM is not an empirical estimator of emotion dynamics from a single video; it is a known-truth data-generating model that lets researchers control latent dimensionality, feedback regime, loading strength, measurement error, and activity pattern before testing whether dynamic methods recover the intended structure.

The paper contributes this model and evaluates its use in three linked steps, shown in Figure 1. First, we specify the LOMM data-generating and measurement components for FER-like video score panels. Second, we compare simulated panels with empirical MAFW-derived FER-like score panels to assess whether the generator reproduces selected temporal and dimensional signatures (Panels A–D). Third, we use the resulting known-truth simulations to evaluate dynamic dimensional recovery with DynEGA, GraphicalVAR, and bounded GIMME under controlled measurement and dynamical conditions (Panels E–H). We then close by discussing when video-derived emotion traces can support behavioral-science claims about latent affective dynamics.

Expressive Affective Dynamics as an Oscillator Measurement Model

Affective-dynamics research treats emotion as a time-varying process rather than as a static score. Across common approaches, the relevant features include average intensity, variability, instability, inertia, transitions between states, and return after perturbation (Hollenstein, 2007; Houben et al., 2015; Kuppens & Verduyn, 2017). For video-derived expressive behavior, the same logic applies to visible affect. A short clip can contain low expressivity, brief expressive episodes, gradual recovery, persistent activation, or alternation

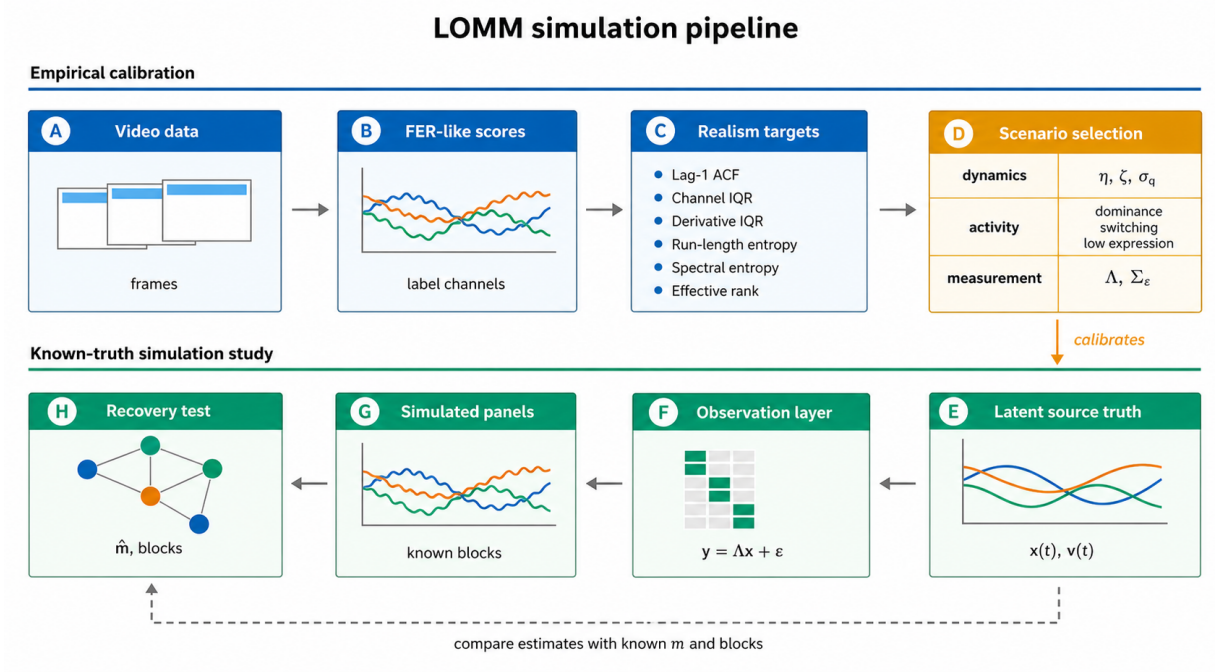


Figure 1: Visual overview of the LOMM simulation pipeline. Panels A–C show the empirical calibration step and realism targets for FER-like video score series. Panel D shows scenario selection through dynamics, activity, and measurement parameters. Panels E–H show the known-truth simulation, read from right to left: latent source-truth trajectories are passed through the observation layer to generate simulated score panels, and recovery tests compare estimated dimensions and blocks with the known latent structure.

between expressive regimes. Modeling expressive affect therefore requires a representation that can describe both the level of a latent expressive process and its tendency to persist, change direction, or return toward equilibrium.

An oscillator model provides a compact way to represent these features once latent variables are understood as theory-guided expressive dimensions that organize FER model outputs rather than as the classifier score labels themselves. Affective structure can be specified in several defensible ways: core-affect and circumplex models organize affect around valence and arousal; PAD models add dominance or control; and Plutchik-style accounts arrange emotion families around opponent tendencies such as joy–sadness, trust–disgust, fear–anger, and anticipation–surprise (Plutchik, 2001; Russell, 1980, 2003; Russell & Mehrabian, 1977). Facial-expression research uses the same distinction between category-like outputs and lower-dimensional structure: AffectNet annotates facial images with both discrete expression categories and continuous valence-arousal ratings, ABAW evaluates valence-arousal estimation alongside expression classification and action-unit detection, and integrated accounts of facial communication treat emotion categories and dimensions as complementary descriptions (Kollias et al., 2023; Mehu & Scherer, 2015; Mollahosseini et al., 2019). For a CLIP-like FER panel, this measurement logic can be made concrete. Frame-level prompt scores for happiness, amusement,

trust, or smiling can be treated as noisy indicators of an approach-affiliation trajectory; anger, fear, disgust, contempt, or anxiety scores can indicate a threat-defense trajectory; and surprise, anticipation, confusion, or attention scores can indicate an orienting-novelty trajectory. LOMM generates the three latent trajectories first, maps them through block loadings into classifier-output score channels, and then lets the recovery method operate only on the simulated score matrix. The known-truth question is therefore whether a method such as DynEGA can recover the intended expressive blocks from CLIP-FER-like outputs whose labels are model scores rather than direct observations of emotion. The model has two integrated parts: a latent oscillator process for expressive dynamics and a measurement model adapted from the direct autoregressive factor score model (H. Golino et al., 2022; Nesselroade et al., 2013).

Following Ollero et al. (2023), LOMM treats each latent expressive dimension as a continuous-time second-order state process. Let $i = 1, \dots, V$ index videos, t index time, and $f = 1, \dots, m$ index the latent expressive dimensions selected for a simulation. For video i, $x_{if}(t)$ is the latent level of dimension f, and $v_{if}(t) = \dot{x}_{if}(t)$ is its instantaneous rate of change. The implemented latent process is

$$\frac{\mathrm{d}}{\mathrm{d}t}\begin{bmatrix} x_{if}(t) \\ v_{if}(t) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ \eta_f & \zeta_f \end{bmatrix}\begin{bmatrix} x_{if}(t) \\ v_{if}(t) \end{bmatrix} + \begin{bmatrix} 0 \\ q_{if}(t) \end{bmatrix}, \qquad q_{if}(t) \sim \mathcal{N}(0, \sigma_{qf}^2), \tag{1}$$

or equivalently, $\ddot{x}_{if}(t) = \eta_f x_{if}(t) + \zeta_f \dot{x}_{if}(t) + q_{if}(t)$. This is an uncoupled multivariate system: η_f, ζ_f, and σ_{qf} may vary by factor, but the model used here does not include off-diagonal dynamic effects in which one latent expressive dimension directly changes the acceleration of another. We therefore interpret η_f as position feedback and ζ_f as velocity feedback. Under this sign convention, negative η_f pulls a displaced latent trajectory back toward its equilibrium level; stronger negative values imply stronger restoring acceleration. Negative ζ_f attenuates velocity over time, $\zeta_f = 0$ removes velocity feedback, and positive ζ_f can sustain or amplify ongoing change over the finite interval being simulated. The innovation $q_{if}(t)$ is a factor-specific shock process that perturbs the latent trajectory. In the simulations reported here, the continuous-time system is advanced with the exact linear transition matrix and initialized with $x_{if}(1) = 0.25$ and $v_{if}(1) = 2$.

For simulation-based evaluation, LOMM defines scenario conditions at two levels. First, the factor-specific parameters η_f, ζ_f, and σ_{qf} set the dynamical regime: negative velocity feedback produces damped return, zero velocity feedback removes damping, and positive velocity feedback can sustain or amplify change over the finite simulated interval. Second, time-varying activity weights determine which latent dimension dominates the generated score series. In single-regime scenarios, one dimension has higher activity across the full sequence; in switching scenarios, the active dimension changes after sampled dwell intervals; and in low-

expression scenarios, all dimensions remain close to background activity. Formally, the trajectory passed to the measurement model is $\tilde{x}_{if}(t) = a_{if}(t)x_{if}(t)$, where $a_{if}(t)$ is a scenario-defined activity weight. This design lets the simulation vary damping, dominance, switching, and volatility while keeping the source-of-truth latent dimensionality and measurement blocks known.

The observable FER-like score vector is generated from the scenario-scaled latent trajectories through a block measurement model. For video i, if each latent dimension has k observed indicators, then $p = mk$ and

$$\mathbf{y}_{it} = {}_i\tilde{\mathbf{x}}_{it} + \varepsilon_{it}, \qquad \varepsilon_{it} \sim \mathcal{N}(\mathbf{0}, {}_i). \tag{2}$$

Each observed score loads on exactly one latent dimension: for indicator j in block B_f, $(\Lambda_i)_{jf} = \lambda + u_{ij}$ with $u_{ij} \sim \text{Unif}(-0.15, 0.15)$, and all off-block loadings are zero. The measurement-error covariance matrix is diagonal, ${}_i = \text{diag}(\sigma^2_{\varepsilon,i1}, \ldots, \sigma^2_{\varepsilon,ip})$, with $\sigma^2_{\varepsilon,ij} = s^2_\varepsilon + r_{ij}$ and $r_{ij} \sim \text{Unif}(-0.01, 0.01)$. In the simulation grid, s_ε is reported as the measurement-error standard deviation, and s^2_ε enters the diagonal covariance matrix. Thus LOMM simulates measured expressive factors and their observed indicator blocks; neutrality is not part of the source-truth structure in the simulation.

For the empirical realism benchmark, this linear block predictor is transformed to the observed score scale with a row-wise softmax link. If $\mathbf{z}_{it} = {}_i\tilde{\mathbf{x}}_{it} + \varepsilon_{it}$ is the pre-link score vector, then the benchmarked observation layer uses

$$y_{itj} = \frac{\exp(z_{itj})}{\sum_{\ell=1}^{p} \exp(z_{it\ell})},$$

equivalent to a softmax temperature of 1.00. No empirical log-mean channel intercept and no additional neutral or background denominator are added in the retained realism benchmark. The softmax link is used to put the empirical realism panels on a bounded, compositional score scale; the known-truth recovery target remains the latent block structure that generated the pre-link score channels.

As a concrete three-factor example, a Plutchik-informed specification could set $m = 3$ for affiliative/approach expression, threat/defensive expression, and orienting/novelty expression (Plutchik, 2001). FER-like indicators related to joy, trust, amusement, or smiling would form one block; indicators related to anger, fear, or disgust would form a second block; and surprise- or anticipation-related indicators would form a third. This specification illustrates the measurement logic of LOMM: the researcher states a theory-guided latent organization, the simulator generates known latent trajectories under that organization, and the classifier-output

score channels are noisy indicators whose loading strength and measurement error are substantive simulation conditions (Kankaanpää et al., 2026). Alternative affective taxonomies can be encoded by changing the latent block map.

The same three-factor example is shown in Figure 2 to make the role of the latent shock term visible.

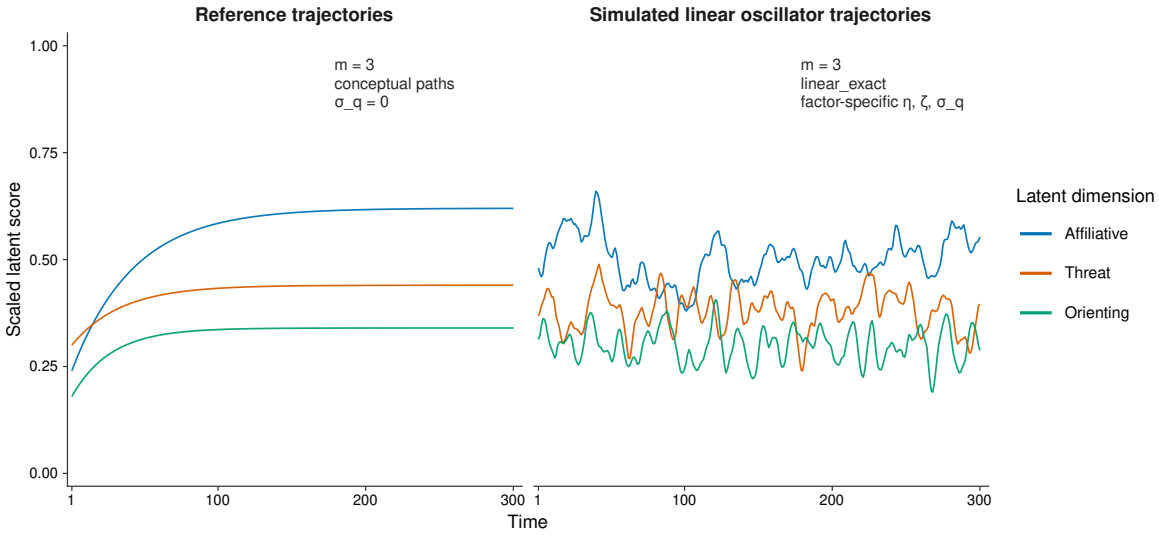


Figure 2: Three-factor LOMM example for Plutchik-informed latent expressive dimensions. The left panel shows smooth reference trajectories without dynamic shocks; the right panel shows trajectories generated by the implemented linear oscillator process with factor-specific position feedback, velocity feedback, and innovation variance.

Figure 2 highlights the role of the innovation variance, σ_q^2, in LOMM. The reference panel shows the conceptual latent ordering without dynamic shocks: affiliative expression rises quickly and settles above the threat and orienting trajectories. The simulated oscillator panel keeps this same structure, but nonzero factor-specific σ_q perturbs each latent process at each time point, producing local peaks, reversals, and irregular fluctuations around the underlying trajectory. After approximately time point 100, affiliative expression (blue) remains the dominant latent tendency, while threat (orange) fluctuates at a moderate level and orienting (green) stays lower with shorter local excursions. Thus, σ_q is the mechanism that turns a smooth conceptual model into latent score series that better resemble the uneven temporal structure of FER-like video data, without adding a neutrality trajectory or cross-factor coupling.

Method

The analyses follow the workflow in Figure 1. We first use empirical video scores to define realism targets for FER-like score series (Panels A–C). We then translate those targets into LOMM scenario conditions

(Panel D) and evaluate dimensional-recovery methods on simulated panels for which the latent dimensional structure and measurement blocks are known (Panels E–H).

Empirical Calibration and Realism Targets

Figure 3 shows the empirical scoring target that motivates the simulation study. The figure uses one clip from MAFW, a dataset designed for dynamic facial expression recognition in the wild (Liu et al., 2022), whereas the realism check uses a 100-clip calibration sample scored with the same zero-shot FER-like procedure. We retrieved frame-level emotion scores with the transforEmotion R package (Tomašević et al., 2026). Each sampled video was decoded into frames, and each frame was evaluated against user-specified emotion labels with the package's `oai-large` vision model, which maps to OpenAI CLIP ViT-Large/14 (`openai/clip-vit-large-patch14`) (Radford et al., 2021). The scoring step returns one frame-level score vector over the prompt-defined channels used here: anger, disgust, fear, happiness, neutral, sadness, surprise, contempt, anxiety, helplessness, and disappointment. Within each clip, short missing runs are interpolated, frames with complete usable scores are retained, and clipped score vectors are renormalized so that frame-level channels can be compared as probability-like outputs. These prompt-specific classifier outputs provide empirical targets for the LOMM observation layer through their temporal persistence, variability, and co-fluctuation.

The selected still frames from the same MAFW clip anchor the empirical example (Panel A). The corresponding score profiles illustrate the measurement problem (Panel B): different frames receive different scores across labels such as helplessness, fear, disgust, sadness, anxiety, and disappointment. The frame-level trajectories show that the highest mean channels change in bursts rather than as smooth curves (Panel C). The single-video DynEGA network applies Dynamic Exploratory Graph Analysis to the same clip's non-neutral score trajectories (Panel D). In this illustration, the method returns three communities among ten score channels, motivating the dimensional-recovery problem later studied under known truth.

The empirical calibration sample was assembled at the clip level from a label-diverse MAFW manifest and then retained after frame-score quality control. The retained panel contains 100 clips spanning the 11 MAFW single-label categories. After interpolation of short missing runs and complete-row filtering, usable clip lengths ranged from 17 to 50 frames, with a median of 36.5 usable frames. These retained frame-score panels define the empirical distributions used for the realism targets.

For the empirical realism check, empirical and simulated panels are summarized at two levels. Observable-channel metrics are computed directly from the FER-like score series within each clip. Community-score

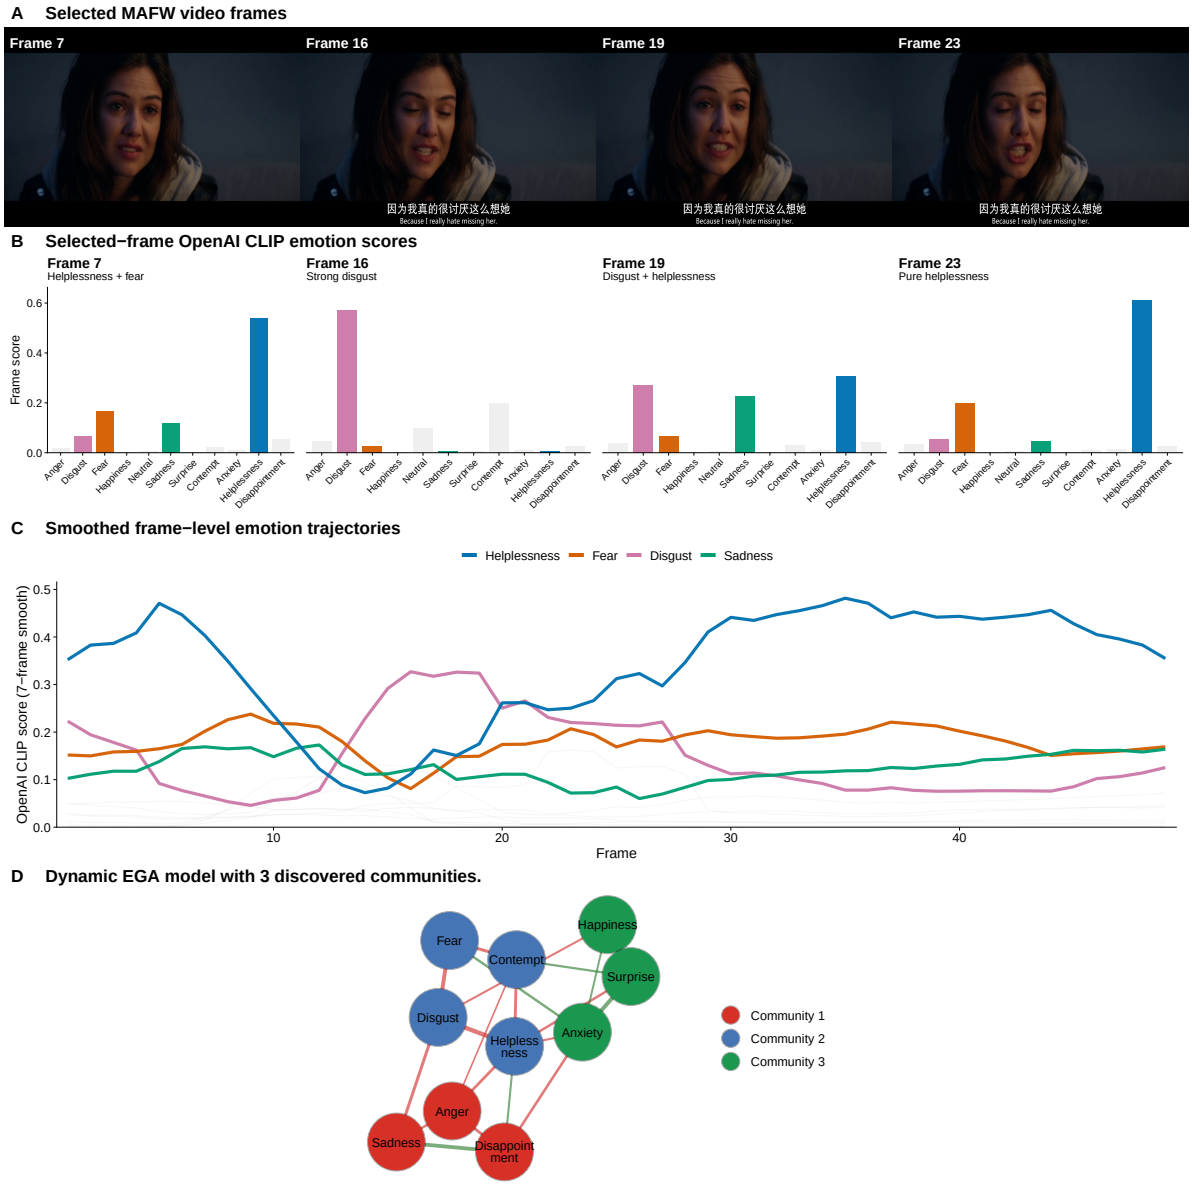


Figure 3: Single-video empirical illustration from the MAFW dataset. Panel A shows selected frames from one MAFW clip; Panel B shows zero-shot FER-like score profiles for the same frames; Panel C shows smoothed frame-level trajectories for the highest mean score channels; Panel D shows a Dynamic EGA network estimated from the same clip's non-neutral score trajectories.

metrics are computed after applying pooled DynEGA to the score panel and extracting score summaries for the recovered communities. This distinction matters because a simulated panel can approximate the observable score channels while still producing a different lower-dimensional geometry, or conversely match the coarse latent organization while retaining imperfect channel-level amplitude behavior.

The primary realism targets include channel-level IQR, derivative IQR, temporal complexity, and dimensional spread. Channel IQR indexes score dispersion; derivative IQR indexes the typical scale of local frame-to-frame movement; lag-1 autocorrelation indexes local persistence; run-length entropy indexes directional-regime complexity; spectral entropy indexes spectral complexity; and effective rank indexes multivariate dimensional spread. For each metric, empirical and simulated distributions are compared using a scaled median difference and a scaled one-dimensional Wasserstein distance, with the larger discrepancy used for classification. We use the same threshold rule across metrics: Pass for scaled discrepancy ≤ 0.33, Partial for scaled discrepancy ≤ 1.00, and Fail for scaled discrepancy > 1.00. These criteria are simulation-construction checks, not external validation of the scoring model or its emotion labels.

Scenario Selection

Scenario selection links the empirical calibration targets to the simulation design. We group the manipulated conditions into three families (Panel D). The dynamics family varies the position feedback, velocity feedback, and innovation variance parameters, η_f, ζ_f, and σ_{qf}. The activity family varies whether one latent expressive dimension dominates across the full sequence, whether dominance switches after sampled dwell intervals, or whether all dimensions remain close to low-expression background activity. The measurement family varies the block loading matrix $_i$ and measurement-error covariance matrix $_i$ that generate the observed FER-like score panels.

The primary emotion-dynamics simulation uses six scenario families crossed with two-, three-, and four-dimensional latent source structures. Four families represent generic expression regimes: low-expression single dominance, high-expression single dominance, high-expression switching, and high-expression fast-switching volatility. Two additional families use the empirical calibration logic more directly, one with ranked persistent dynamics and one with broader switching geometry. In all cases, the number of observed channels is $p = mk$, where m is the number of latent dimensions and k is the number of observed indicators per dimension. The primary design uses $N = 50$ or 100 videos, $T = 100$ time points, $k = 4$ observed indicators per latent dimension, measurement-error standard deviations of 0.125, 0.25, and 0.50, loadings of 0.4, 0.6, and 0.8, and three replications per structural cell. Two sensitivity conditions hold the signal settings fixed while increasing the number of indicators to $k = 6$ or the sequence length to $T = 300$.

For each simulated video, LOMM first generates latent trajectories from the scenario-specific dynamic parameters. The activity schedule then scales those trajectories: non-switching scenarios keep one active dimension for the full sequence, rotating the dominant dimension across videos, whereas switching scenarios sample dwell intervals and change the active dimension over time. The measurement layer is applied to the activity-scaled latent trajectories. Measurement error is reported as a standard deviation, with the corresponding variance entering the diagonal measurement-error covariance before small item-level variance jitter is added. The simulation design evaluates dimensional recovery under controlled combinations of dynamics, activity regimes, and measurement quality.

Dimensional-Recovery Simulation Study

The simulation begins with latent trajectories and the known number of latent expressive dimensions (Panel E). The LOMM observation layer maps each observed score channel to one latent dimension through the block loading matrix and adds item-level measurement error (Panel F), producing a simulated FER-like score panel (Panel G). The final step is the dimensional-recovery test, where the estimated number of dimensions and item memberships are compared against the known latent structure (Panel H).

We compare Dynamic Exploratory Graph Analysis [DynEGA; H. Golino et al. (2022)] with GraphicalVAR (Epskamp et al., 2018; Epskamp, 2024) and GIMME (Gates & Molenaar, 2012; Lane et al., 2019) because they represent adjacent but distinct approaches to intensive longitudinal structure. DynEGA is the focal dimensional-recovery method: it estimates groups of variables whose changes are coordinated over time. GraphicalVAR is a dynamic-network comparator that separates lagged temporal relations from contemporaneous partial associations. GIMME is an idiographic time-series framework that searches for relations that are common across individuals, shared by subgroups, or unique to individuals. In the present simulation, DynEGA tests whether a method designed for dimensional recovery can recover the known score blocks directly, GraphicalVAR tests whether a dynamic-network representation yields a similar partition after community detection, and GIMME tests whether a more flexible person-specific time-series framework is feasible under the same measurement conditions.

DynEGA combines time-delay embedding and derivative estimation from dynamical systems with exploratory graph analysis for dimensionality assessment (Boker et al., 2010; H. Golino et al., 2020; H. F. Golino & Epskamp, 2017). We apply DynEGA to the stacked video panel with video identity supplied as the grouping variable. The embedding dimension is selected by minimizing the Total Entropy Fit Index (TEFI) over candidate values (H. Golino et al., 2021), using $n_{\text{embed}} \in \{5, 10, 15, \lfloor T/4 \rfloor, \lfloor T/2 \rfloor\}$ where admissible. The estimated DynEGA population network is then partitioned into communities, which are interpreted

as recovered groups of score channels. The expanded DynEGA grid uses Louvain community detection (Blondel et al., 2008) with the expand rule for unidimensional solutions. In the bounded all-method comparator, community detection is analyzed as a factor with Louvain, Leiden (Traag et al., 2019), and Walktrap (Pons & Latapy, 2005), holding the unidimensional rule at expand.

GraphicalVAR estimates a graphical vector autoregressive model for intensive longitudinal data (Epskamp et al., 2018; Epskamp, 2024). In this framework, the temporal network contains directed lag-1 regression parameters, whereas the contemporaneous network contains partial correlations among residuals at the same time point after lagged effects have been taken into account. In multilevel applications, the model can also distinguish within-person temporal and contemporaneous structure from between-person associations among stationary means. We fit GraphicalVAR to the stacked simulated score panel and use the estimated contemporaneous partial-correlation matrix as the structural matrix for dimensional partitioning, because the LOMM recovery target is an undirected grouping of score channels rather than a directed lagged-effect graph. The diagonal is ignored and the resulting weighted matrix is treated as an undirected graph. Positive edge weights define the graph, with absolute weights used only when no positive edges are available. Louvain, Leiden, and Walktrap are applied to the same fitted contemporaneous structure and are analyzed as the community-detection factor.

GIMME is an automated person-specific time-series modeling approach built around unified structural equation modeling and vector autoregressive representations (Gates & Molenaar, 2012; Lane et al., 2019). The algorithm searches for lagged and contemporaneous paths that improve fit across the sample, retains paths that are common to many individuals, and then allows subgroup-shared or individual-specific paths so that heterogeneous time-series structures do not have to be forced into one pooled model. In this simulation, each video is treated as the idiographic time series supplied to GIMME. We fit the VAR specification to the video-level score panels, extract the contemporaneous block from the person-specific Ψ matrices, and average those blocks across videos to obtain a common structural matrix for dimensional partitioning. The diagonal is ignored and the averaged matrix is treated as an undirected graph. As with GraphicalVAR, positive edge weights define the graph, with absolute weights used only when no positive edges are available. Louvain, Leiden, and Walktrap are applied to the same fitted GIMME structure and are analyzed as the community-detection factor.

All three methods are evaluated as dimensional-recovery estimators under known truth. DynEGA is evaluated on an expanded grid for scalability and on the bounded all-method comparator; GraphicalVAR and GIMME are evaluated on the bounded comparator because complete factorial estimation is computationally unstable in this setting. Community-detection algorithm is included as a diagnostic factor in the bounded

comparator rather than as the primary organizing variable. Performance is summarized using dimensional hit rate, assignment accuracy for observed score channels, matched-score correlation, matched-score normalized root mean squared error (NRMSE), runtime, and failure or timeout accounting. Dimensional hit rate is the proportion of method-readout rows in which the method returns the correct number of latent dimensions. Assignment accuracy asks whether observed score channels are assigned to the correct latent blocks after optimal matching of estimated communities to true blocks. For score-based metrics, observed variables are standardized within the full panel and community scores are computed from the first principal component of each estimated block, with row means used when a principal-component score cannot be estimated. Matched-score correlation and matched-score NRMSE then compare these recovered score summaries with the known latent trajectories after sign-invariant matching between estimated communities and source dimensions. Runtime is recorded for each model fit. Non-converged fits, post-estimation extraction errors, and predefined timeouts remain in the denominators. For GraphicalVAR and GIMME, the same fitted structure is partitioned by multiple community-detection algorithms, so feasibility is also reported as unique structure-fit status counts.

The final simulation uses a split design. The expanded DynEGA grid varies the number of videos, series length, latent dimensionality, indicators per dimension, loading strength, and measurement-error standard deviation across 1,296 structural cells with 400 replications, yielding 518,400 DynEGA Louvain/expand method rows. The all-method comparator uses the bounded NVAR = 4 design with two- and three-dimensional source structures, two sample sizes, TimeP = 100, three loading strengths, three measurement-error standard deviations, and all six scenario families. This bounded comparator contains 216 structural cells with 100 replications. Louvain, Leiden, and Walktrap are included as community-detection levels, yielding 64,800 method-readout rows per method in the final comparator. Timeouts, failed fits, and post-estimation extraction errors are retained in the reporting denominator because they are part of method behavior in applied simulation workflows.

Results

Results are organized around four questions that follow the workflow in Figure 1. We first ask whether the calibrated LOMM generator produces FER-like score panels with empirical temporal and dimensional signatures. We then evaluate whether dimensional-recovery methods recover the known latent structure, examine how recovery changes with measurement quality, and report the computational feasibility of the compared methods.

Empirical Realism Results

Before evaluating dimensional-recovery methods, we asked whether the calibrated LOMM generator produced FER-like score panels that were sufficiently close to empirical panels for simulation construction. Figure 4 summarizes this check at two readout levels: the observed score channels and the lower-dimensional community scores extracted from the pooled DynEGA solution. The comparison evaluates selected distributional, temporal, spectral, and dimensional signatures of the empirical MAFW score panels, providing a simulation-construction check on the simulated panels.

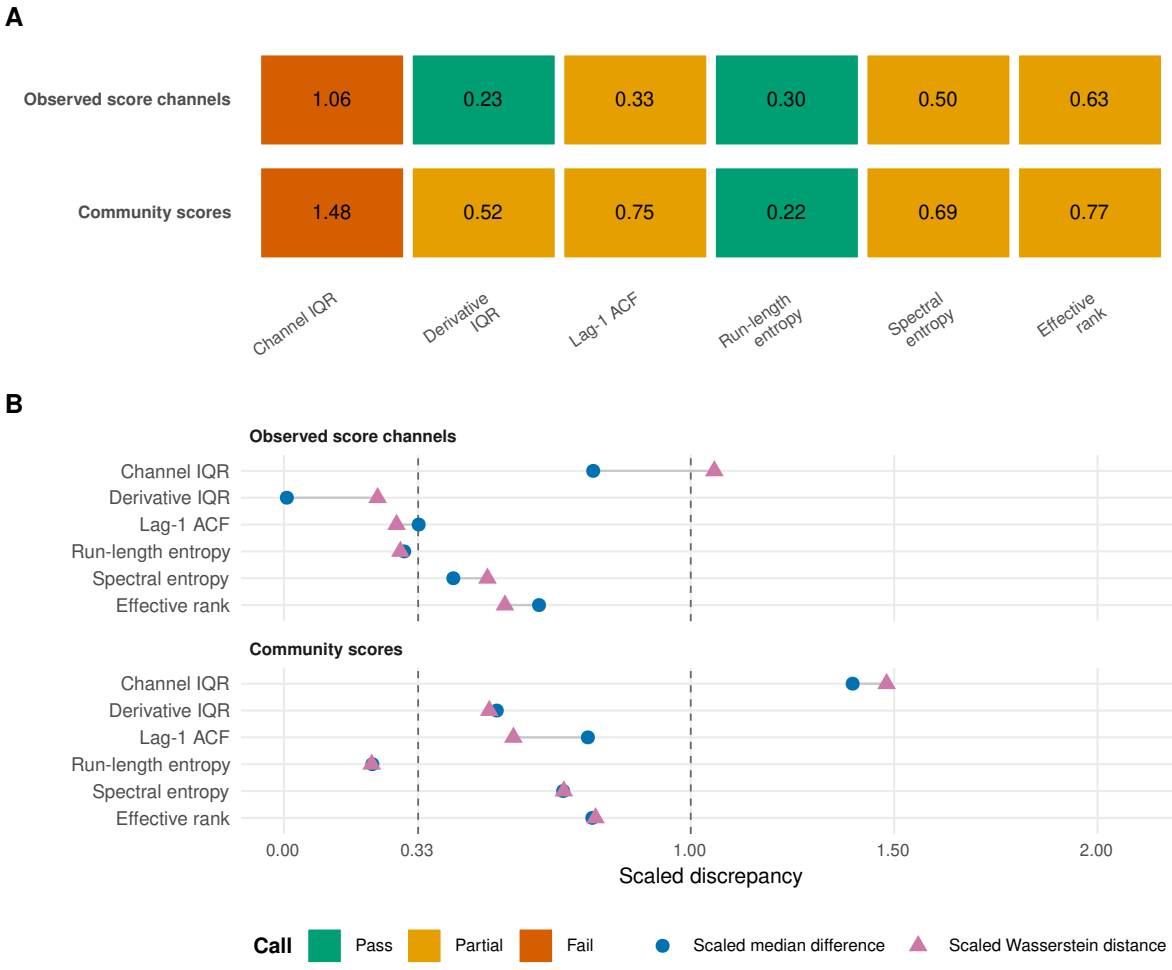


Figure 4: Empirical realism check for LOMM-generated FER-like panels. Panel A shows pass, partial, and fail classifications for the six realism targets at two readout levels: observed score channels and DynEGA community scores. Values are the larger of the scaled median difference and scaled Wasserstein distance. Panel B shows the two discrepancy components separately; dashed reference lines mark the pass threshold of 0.33 and fail threshold of 1.00.

At the observed-channel level, the softmax-linked generator passed the derivative-IQR and run-length entropy checks. Lag-1 autocorrelation, spectral entropy, and effective rank were partial matches, indicating that the

simulated panels were close to the empirical panels in local persistence, spectral complexity, and multivariate dimensional spread but not within the strictest threshold. The only failed observed-channel check was channel IQR, showing that the simulated observed channels remained more dispersed than the empirical score channels.

The community-score readout showed the same broad pattern with stronger amplitude mismatch. Run-length entropy passed, derivative IQR, lag-1 autocorrelation, spectral entropy, and effective rank were partial matches, and channel IQR failed. The retained Louvain community readout returned three communities with the same 6-3-2 channel-size pattern in the empirical and simulated panels, so the coarse dimensional organization matched in this readout even though the marginal amplitude of the community scores was too large. Taken together, these results support using the softmax-linked generator for dimensional-recovery simulation while keeping the amplitude and dispersion mismatch explicit as a limitation of the current calibration.

Simulation Results

The final recovery analysis separates scalability, direct method comparison, and factor diagnostics. DynEGA was first evaluated on the expanded Louvain/expand grid, which contains 1,296 structural cells and 518,400 method rows. The direct method comparison then used the bounded NVAR = 4 design shared by DynEGA, GraphicalVAR, and GIMME. This bounded comparator contains 216 structural cells and 100 replications per cell. Louvain, Leiden, and Walktrap were included as community-detection readouts, yielding 64,800 method-readout rows per method, but the main comparison aggregates across these readouts and then evaluates community detection in the factor screen. Table 1 summarizes dimensional recovery, block assignment, score recovery, status counts, and runtime by method.

Table 1: Overall method behavior in the final bounded NVAR = 4 comparator, aggregated across Louvain, Leiden, and Walktrap readouts. Hit rate is recovery of the known number of latent dimensions using all method-readout rows as the denominator; Assign. is optimal block-assignment accuracy; Score r and Score NRMSE compare matched recovered score summaries with the known latent trajectories.

Method	Rows	ok	Timeouts	Errors	Hit (%)	Assign. (%)	Score r	NRMSE	Time (s)
DynEGA	64800	64800	0	0	95.3	96.0	0.725	0.649	2.61
GraphicalVAR	64800	62616	2184	0	80.3	90.4	0.717	0.659	86.10
GIMME	64800	45882	12342	6576	25.9	67.2	0.604	0.842	280.50

DynEGA was the only method evaluated on the expanded grid and completed all 518,400 Louvain/expand rows. On that grid, it recovered the correct number of dimensions in 93.7% of rows and assigned observed channels to the correct blocks with 95.2% accuracy. The bounded comparator in Table 1 then places DynEGA,

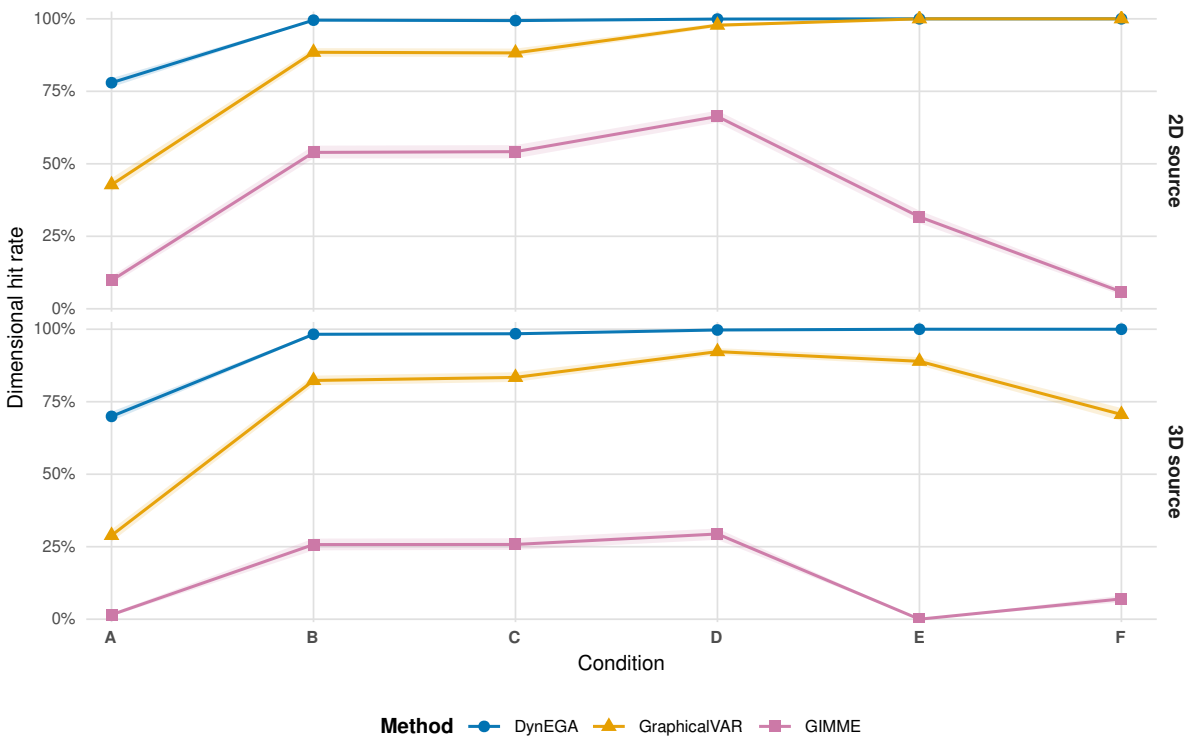


Figure 5: Dimensional hit rate across the final bounded NVAR = 4 method-comparison conditions. Lines show methods and shaded bands show cluster-bootstrap 95% confidence intervals, resampling generated panels while retaining their community-detection readouts. The condition codes are A = low-expression single-dominance, B = high-expression single-dominance, C = high-expression switching, D = high-expression rapid switching, E = ranked-persistence, and F = broad switching-geometry. Hit rates are status-inclusive and aggregated across Louvain, Leiden, and Walktrap readouts.

GraphicalVAR, and GIMME on the same NVAR = 4 design. This comparison keeps unsuccessful fits in the denominator, so the reported hit rates combine structural recovery and estimator feasibility. Aggregating across the three community-detection readouts, DynEGA recovered the correct number of dimensions in 95.3% of bounded rows, GraphicalVAR in 80.3%, and GIMME in 25.9%. Assignment accuracy followed the same ordering, whereas conditional score summaries were interpreted separately because they are only meaningful after structural recovery.

Figure 5 shows the same bounded-comparator results graphically. DynEGA retained the strongest recovery profile overall, but the scenario-family pattern is important: GraphicalVAR recovered the known dimensionality in many OpenAI-style conditions when the structure fit completed, while recovery dropped in low-signal old-style scenarios and some OpenAI-style three-dimensional structure fits reached the timeout boundary. GIMME showed the lowest status-inclusive recovery because many OpenAI-style fits timed out or produced no usable contemporaneous structure for partitioning. The main simulation result is therefore a combination of target alignment, scenario sensitivity, and feasibility: DynEGA recovered the known block structure most consistently, GraphicalVAR was condition-sensitive but often recoverable, and GIMME was strongly feasibility-limited in this FER-like bounded design.

The descriptive factor screen for the bounded comparator treats method, community detection, scenario family, source dimensionality, loading strength, and measurement-error standard deviation as predictors of recovery and score-quality summaries. For dimensional hit rate, method and scenario family accounted for the largest differences ($\eta_p^2 = .404$ and .241), measurement-error standard deviation and loading strength had smaller but meaningful effects ($\eta_p^2 = .146$ and .036), and community detection accounted for essentially no variation ($\eta_p^2 < .001$). The same pattern held for estimator completion, block assignment, score correlation, and score NRMSE: community-detection effects rounded to zero, whereas method, scenario family, and measurement quality accounted for the substantive variation. Runtime was the only outcome with a nonzero community-detection contribution ($\eta_p^2 = .049$), still much smaller than method ($\eta_p^2 = .944$) and scenario family ($\eta_p^2 = .294$). Because scenario differences are shown in Figure 5, the next comparison focuses on the two measurement-quality factors that directly alter the observed FER-like panels: loading strength and measurement error.

Figure 6 shows how method differences changed across measurement quality. DynEGA was most stable across the loading-by-error grid, although low-signal scenarios still reduced recovery. GraphicalVAR was more sensitive to the measurement layer: its hit rate was weakest when loadings were low and measurement error was high, and improved when observed channels carried a stronger signal from their source dimensions. GIMME remained lower on status-inclusive hit rate across most strata because errors and timeouts occurred

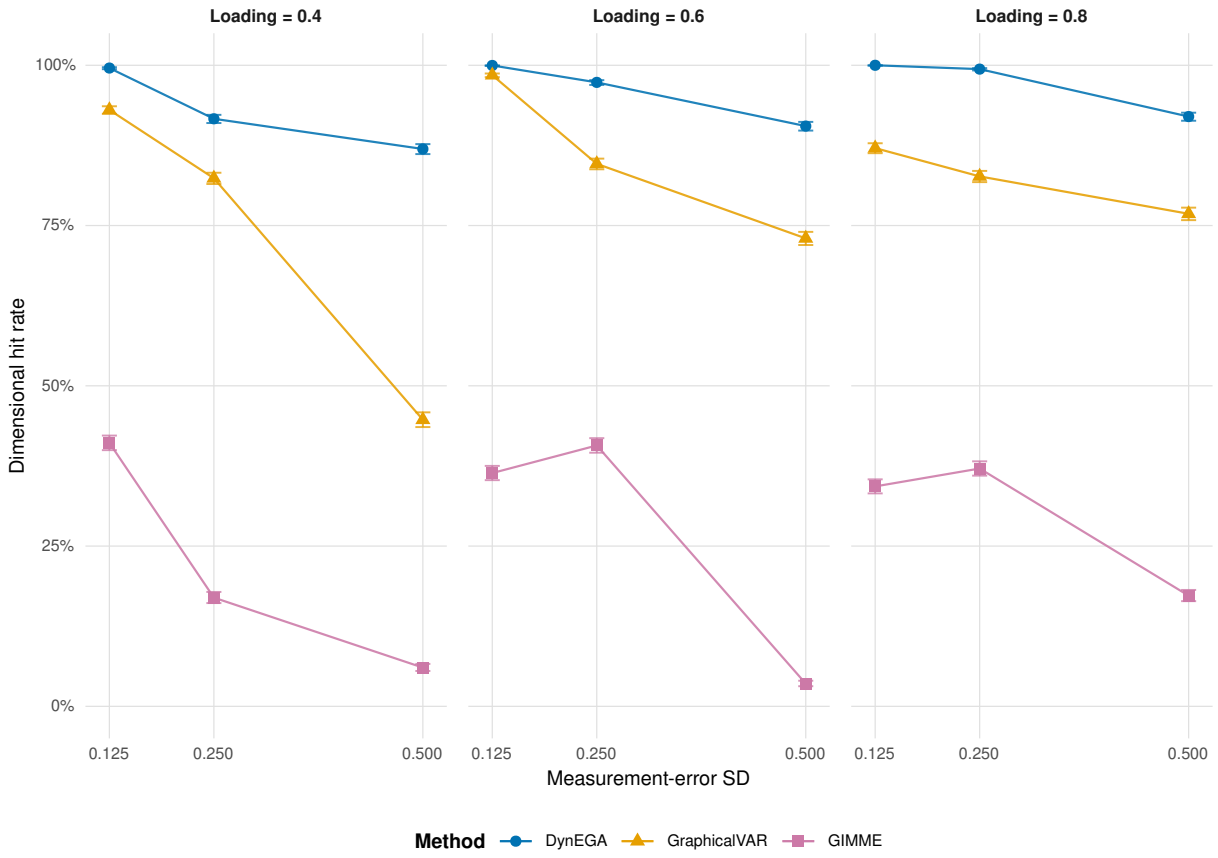


Figure 6: Dimensional hit rate stratified by measurement-error standard deviation and loading strength in the final bounded NVAR = 4 comparator. Columns show loading strength; lines and symbols show methods. Error bars are Wilson 95% confidence intervals computed from all method-readout rows in each stratum, aggregated across community-detection readouts.

before dimensional recovery could be evaluated.

Table 2: Conditional matched-score correlations after correct dimensional recovery in the final bounded NVAR = 4 comparator. Correlations use only rows in which the corresponding method readout recovered the correct number of dimensions. The conditional r range and widest CI are computed across the loading-by-measurement-error strata shown in Figure 6.

Method	Correct	Hit (%)	Mean r	r range	Stratum n	Missing
DynEGA	61732/64800	95.3%	0.764	0.585-0.877	6260-7200	0
GraphicalVAR	52041/64800	80.3%	0.802	0.726-0.858	3219-7089	0
GIMME	16804/64800	25.9%	0.797	0.759-0.951	255-2959	0

Table 2 shows why matched-score correlations should be interpreted as a conditional diagnostic rather than a primary method-ranking criterion. Conditional correlations use only readout rows in which the correct number of dimensions was recovered, so the denominators differ sharply across methods. DynEGA contributed 61,732 correctly recovered rows, GraphicalVAR contributed 52,041, and GIMME contributed 16,804. Within those successful rows, mean conditional correlations were high for all three methods, but the wider GIMME intervals and smaller stratum counts show that this conditional score result does not remove the feasibility and structural-recovery differences shown in Table 1 and Figure 5.

Computational Feasibility

Computational feasibility was evaluated separately from recovery accuracy because the methods differed in both recovery behavior and scalability across the simulation design. Figure 7 summarizes successful-fit runtime in the final bounded comparator, and the text below accounts for full-grid completion, timeouts, and post-estimation errors. This distinction is important for interpreting the comparator results: GraphicalVAR and GIMME are informative bounded comparisons with more limited scalability across the full LOMM grid.

DynEGA was substantially faster than the two model-based comparators in the bounded design. Median successful-fit runtime was 2.61 s for DynEGA, compared with 86.1 s for GraphicalVAR and 280.5 s for GIMME. The expanded DynEGA Louvain/expand grid also completed all 518,400 method rows without errors or timeouts, supporting its use for the broader scalability screen. GraphicalVAR completed 20,872 of 21,600 bounded NVAR = 4 structure fits and timed out on 728 fits, concentrated in OpenAI-style three-dimensional scenarios. GIMME completed 15,294 of 21,600 bounded structure fits, timed out on 4,114 fits, and returned post-estimation errors for 2,192 fits. Because GraphicalVAR and GIMME apply three community-detection algorithms to each fitted structure, these feasibility counts are unique structure-fit counts rather than repeated CD readout rows. The final simulation results therefore support two linked conclusions: DynEGA had the strongest dimensional-recovery profile in the bounded comparator and was

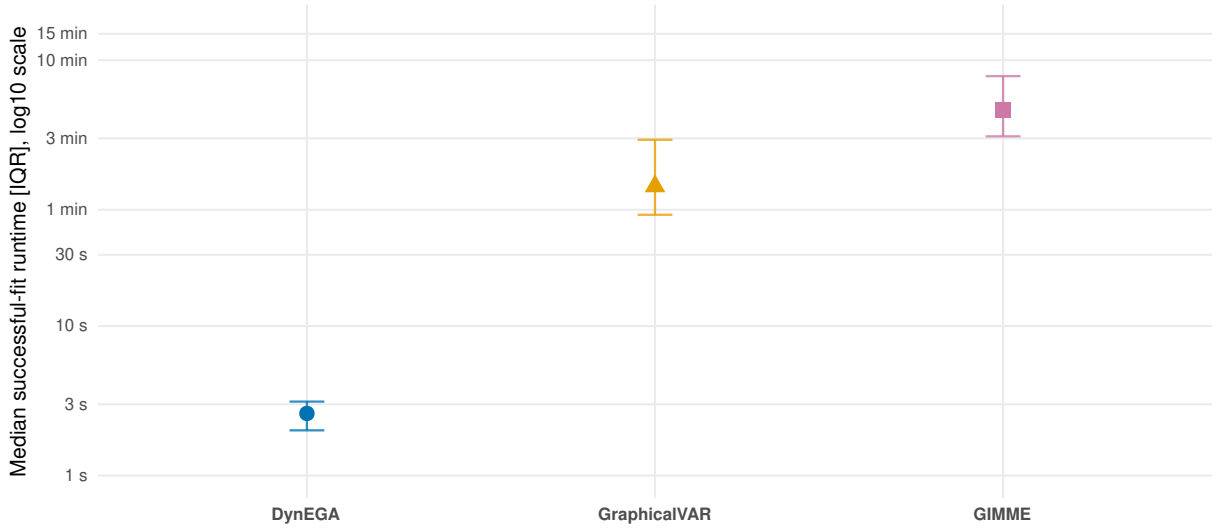


Figure 7: Runtime of the compared methods in the final bounded NVAR = 4 comparator. The y-axis is shown on a log10 time scale. Points show median runtime per successful method-readout row and intervals show the interquartile range, aggregated across community-detection readouts. For GraphicalVAR and GIMME, community-detection readouts share the same fitted structure within a panel.

the only method light enough to cover the expanded LOMM grid; GraphicalVAR and GIMME provide informative bounded comparators whose interpretation must include timeout and extraction-error behavior.

Discussion

This study introduced LOMM as a known-truth framework for studying dimensional recovery in FER-like video score series. The strongest conceptual result is that LOMM turns dynamic dimensional recovery into a controlled perturbation problem. In the population design, m latent oscillator processes generate m measured score blocks. Within a block, observed channels share the same latent level and rate of change up to loading differences, activity weights, and measurement error. The recovery task is therefore to recover m rank-one dynamic subspaces from noisy, activity-gated, block-measured oscillator traces.

This framing clarifies why the method comparison is mainly about target alignment and scalability. DynEGA is designed to estimate groups of variables whose changes are coordinated over time, and this target is close to the population object generated by LOMM. In the ideal case, each score block forms a coherent dynamic subspace; finite samples, measurement error, weak loadings, activity switches, and oscillator amplitude differences perturb that block structure. DynEGA's high dimensional hit rate, high block-assignment accuracy, and complete expanded-grid coverage show that the estimated dynamic communities usually remained close enough to the known LOMM blocks for recovery, while the low-signal scenarios identify the boundary of

that performance. The community-detection diagnostic did not materially change this conclusion because its effect on recovery and score-quality outcomes was negligible relative to method, scenario family, and measurement quality.

GraphicalVAR is a useful comparator because it estimates dynamic network structure, but its structural matrix is less directly aligned with a reflective block measurement model. The contemporaneous partial-correlation representation conditions away shared variance among indicators. Under a common-cause block structure, this can attenuate the within-block edges that define the latent measurement block, even when the marginal block signal is strong. This population alignment explains why finite-sample noise, lower loadings, higher measurement error, and temporal-model estimation can more easily disrupt the recovered partition. In the final bounded results, community-detection choice made very little difference once the GraphicalVAR structure fit had been obtained; the main limits were scenario sensitivity and OpenAI-style three-dimensional timeouts.

GIMME addresses a broader problem than the one LOMM makes primary. It can represent group-level, subgroup-level, and individual-level time-series structure, which makes it attractive for idiographic dynamic modeling. In this design, that flexibility came with weak status-inclusive dimensional hit rates, long runtimes, frequent timeouts, and post-estimation errors in the OpenAI-style panels. GIMME therefore contributes mainly as a feasibility-stress comparator for this simulation rather than as a scalable full-grid peer.

The score-recovery results reinforce the same interpretation. Once a method recovered the correct block structure, recovered score summaries often tracked the known latent trajectories well. This is expected because a correctly recovered score block has an oracle-like reliability determined mainly by loading strength, measurement error, number of indicators, and latent amplitude. The larger methodological difference was the probability of reaching the correct structure in the first place and producing a usable fitted structure within the time limit. In this sense, method performance in LOMM decomposes into structural recovery probability, conditional score reliability, and feasibility, with DynEGA's advantage coming primarily from the first and third components.

The empirical realism check sets the boundary on this interpretation. The softmax-linked simulated panels matched the empirical FER-like score panels well for local movement scale and directional-regime texture, and they partially matched persistence, spectral complexity, and dimensional spread. Channel IQR failed at both the observed-channel and community-score readouts, showing that the current calibration still produces larger marginal amplitude dispersion than the empirical panels. This mismatch is informative rather than disqualifying: LOMM can reproduce selected dynamic and dimensional signatures while retaining a narrower

score-range and occupancy limitation in the bounded, probability-like observation layer.

Several limitations follow directly from this design. FER-like scores remain model-derived observations of visible expression, so the realism check evaluates selected statistical signatures rather than the validity of emotion labels or exact empirical trajectories. The present LOMM implementation uses independent oscillator dimensions, activity-gated block loadings, and diagonal measurement error; it does not claim to capture the full geometry of real affective systems. Future work should test richer measurement structures, bounded or compositional observation layers, additional empirical scoring systems, and empirical applications where known-truth simulation results guide substantive validation of video-derived affective dynamics.

Declarations

Funding

The authors received no specific funding for this work.

Competing interests

The authors declare that they have no competing interests.

Ethics approval

This study used simulations and secondary, publicly described benchmark materials. It did not involve new recruitment, intervention, or direct collection of human-participant data by the authors.

Consent to participate

Not applicable.

Consent for publication

Not applicable.

Availability of data and materials

The manuscript analyses use simulation outputs and derived FER-like score summaries generated for this study. The version-controlled manuscript, analysis code, figure-generation scripts, simulation scripts, and reproducibility materials are available at https://github.com/atomashevic/dlo-rr. Raw MAFW video materials are third-party dataset materials; access to the original videos is governed by the MAFW dataset providers. A permanent OSF archival record will be added before final journal upload.

Code availability

All code needed to reproduce the manuscript figures, tables, and simulation summaries is available in the project repository at https://github.com/atomashevic/dlo-rr. The repository includes the R scripts used for simulation, empirical realism summaries, method-comparison summaries, and manuscript rendering.

Authors' contributions

Aleksandar Tomasevic contributed conceptualization, formal analysis, methodology, resources, software, validation, visualization, writing the original draft, and reviewing and editing the manuscript. Hudson Golino contributed formal analysis, methodology, resources, software, validation, visualization, writing the original draft, and reviewing and editing the manuscript. Alexander P. Christensen contributed methodology, software, writing the original draft, and reviewing and editing the manuscript.

Open Practices Statement

The study was not preregistered. Analysis code, simulation code, manuscript source files, generated figure and table inputs, and derived reproducibility materials are available at https://github.com/atomashevic/dlo-rr. Raw MAFW video data are third-party materials governed by the dataset providers. A permanent OSF archive will be linked before final journal upload.

References

Arriaga, O., Valdenegro-Toro, M., & Plöger, P. (2017). *Real-time Convolutional Neural Networks for Emotion and Gender Classification* (arXiv:1710.07557). arXiv. https://doi.org/10.48550/arXiv.1710.07557

Barrett, L. F., Adolphs, R., Marsella, S., Martinez, A. M., & Pollak, S. D. (2019). Emotional Expressions Reconsidered: Challenges to Inferring Emotion From Human Facial Movements. *Psychological Science in the Public Interest*, *20*(1), 1–68. https://doi.org/10.1177/1529100619832930

Blondel, V. D., Guillaume, J.-L., Lambiotte, R., & Lefebvre, E. (2008). Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, *2008*(10), P10008. https://doi.org/10.1088/1742-5468/2008/10/P10008

Boker, S. M. (2015). Adaptive equilibrium regulation: A balancing act in two timescales. *Journal for Person-Oriented Research*, *1*(1-2), 99–109. https://doi.org/10.17505/jpor.2015.10

Boker, S. M., Deboek, P. R., Edler, C., & Keel, P. (2010). Generalized local linear approximation of derivatives from time series. In S. M. Chow, E. Ferrer, & F. Hsieh (Eds.), *The notre dame series on quantitative methodology. Statistical methods for modeling human dynamics: An interdisciplinary dialogue* (pp. 161–178). Routledge/Taylor & Francis Group.

Boker, S. M., & Nesselroade, J. R. (2002). A Method for Modeling the Intrinsic Dynamics of Intraindividual Variability: Recovering the Parameters of Simulated Oscillators in Multi-Wave Panel Data. *Multivariate Behavioral Research*, *37*(1), 127–160. https://doi.org/10.1207/S15327906MBR3701_06

Epskamp, S. (2024). *graphicalVAR: Graphical VAR for experience sampling data*. https://CRAN.R-project.org/package=graphicalVAR

Epskamp, S., Waldorp, L. J., Mõttus, R., & Borsboom, D. (2018). The gaussian graphical model in cross-sectional and time-series data. *Multivariate Behavioral Research*, *53*(4), 453–480. https://doi.org/10.1080/00273171.2018.1454823

Ezerceli, Ö., & Eskil, M. T. (2022). Convolutional Neural Network (CNN) Algorithm Based Facial Emotion Recognition (FER) System for FER-2013 Dataset. *2022 International Conference on Electrical, Computer, Communications and Mechatronics Engineering (ICECCME)*, 1–6. https://doi.org/10.1109/ICECCME55909.2022.9988371

Foteinopoulou, N. M., & Patras, I. (2023). *EmoCLIP: A Vision-Language Method for Zero-Shot Video Facial Expression Recognition* (arXiv:2310.16640). arXiv. https://doi.org/10.48550/arXiv.2310.16640

Gates, K. M., & Molenaar, P. C. M. (2012). Group search algorithm recovers effective connectivity maps for individuals in homogeneous and heterogeneous samples. *NeuroImage*, *63*(1), 310–319. https://doi.org/10.1016/j.neuroimage.2012.06.026

Golino, H. F., & Epskamp, S. (2017). Exploratory graph analysis: A new approach for estimating the number of dimensions in psychological research. *PloS One*, *12*(6), e0174035. https://doi.org/10.1371/journal.pone.0174035

Golino, H., Christensen, A. P., Moulder, R., Kim, S., & Boker, S. M. (2022). Modeling latent topics in social media using dynamic exploratory graph analysis: The case of the right-wing and left-wing trolls in the 2016 US elections. *Psychometrika*, *87*(1), 156–187.

Golino, H., Moulder, R., Shi, D., Christensen, A. P., Garrido, L. E., Nieto, M. D., Nesselroade, J., Sadana, R., Thiyagarajan, J. A., & Boker, S. M. (2021). Entropy fit indices: New fit measures for assessing the structure and dimensionality of multiple latent variables. *Multivariate Behavioral Research*, *56*(6), 874–902. https://doi.org/10.1080/00273171.2020.1779642

Golino, H., Shi, D., Garrido, L. E., Christensen, A. P., Nieto, M. D., Sadana, R., Thiyagarajan, J. A., & Martinez-Molina, A. (2020). Investigating the performance of exploratory graph analysis and traditional techniques to identify the number of latent factors: A simulation and tutorial. *Psychological Methods*, *25*(3), 292–230. https://doi.org/10.1037/met0000255

Goodfellow, I. J., Erhan, D., Carrier, P. L., Courville, A., Mirza, M., Hamner, B., Cukierski, W., Tang, Y., Thaler, D., Lee, D.-H., Zhou, Y., Ramaiah, C., Feng, F., Li, R., Wang, X., Athanasakis, D., Shawe-Taylor, J., Milakov, M., Park, J., … Bengio, Y. (2013). *Challenges in Representation Learning: A report on three machine learning contests* (arXiv:1307.0414). arXiv. https://doi.org/10.48550/arXiv.1307.0414

Halamová, J., Kanovský, M., Brockington, G., & Strnádelová, B. (2023). Automated facial expression

analysis of participants self-criticising via the two-chair technique: Exploring facial behavioral markers of self-criticism. *Frontiers in Psychology*, *14*. https://doi.org/10.3389/fpsyg.2023.1138916

Hoemann, K., Khan, Z., Feldman, M. J., Nielson, C., Devlin, M., Dy, J., Barrett, L. F., Wormwood, J. B., & Quigley, K. S. (2020). Context-aware experience sampling reveals the scale of variation in affective experience. *Scientific Reports*, *10*(1), 12459. https://doi.org/10.1038/s41598-020-69180-y

Holiday, S., Hayes, J. L., Park, H., Lyu, Y., & Zhou, Y. (2023). A multimodal emotion perspective on social media influencer marketing: The effectiveness of influencer emotions, network size, and branding on consumer brand engagement using facial expression and linguistic analysis. *Journal of Interactive Marketing*, *58*(4), 414–439. https://doi.org/10.1177/10949968231171104

Hollenstein, T. (2007). State space grids: Analyzing dynamics across development. *International Journal of Behavioral Development*, *31*(4), 384–396. https://doi.org/10.1177/0165025407077765

Houben, M., Van Den Noortgate, W., & Kuppens, P. (2015). The relation between short-term emotion dynamics and psychological well-being: A meta-analysis. *Psychological Bulletin*, *141*(4), 901–930. https://doi.org/10.1037/a0038822

Kankaanpää, R., Ron, J. de, Hoekstra, R. H. A., & Bork, R. van. (2026). Comparing multiple-indicator approaches to account for measurement error in dynamic networks. *Cognitive Therapy and Research*. https://doi.org/10.1007/s10608-026-10719-0

Kollias, D., Tzirakis, P., Baird, A., Cowen, A., & Zafeiriou, S. (2023). ABAW: Valence-arousal estimation, expression recognition, action unit detection, and emotional reaction intensity estimation challenges. *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 5889–5898. https://doi.org/10.1109/CVPRW59228.2023.00626

Kuppens, P., & Verduyn, P. (2017). Emotion dynamics. *Current Opinion in Psychology*, *17*, 22–26. https://doi.org/10.1016/j.copsyc.2017.06.004

Lane, S. T., Gates, K. M., Pike, H. K., Beltz, A. M., & Wright, A. G. C. (2019). Uncovering general, shared, and unique temporal patterns in ambulatory assessment data. *Psychological Methods*, *24*(1), 54–69. https://doi.org/10.1037/met0000192

Li, H., Niu, H., Zhu, Z., & Zhao, F. (2023). *CLIPER: A Unified Vision-Language Framework for In-the-Wild Facial Expression Recognition* (arXiv:2303.00193). arXiv. https://doi.org/10.48550/arXiv.2303.00193

Li, Q., Liu, Z., Zhang, Z., Wang, Q., & Ma, M. (2024). Decoding group emotional dynamics in a web-based collaborative environment: A novel framework utilizing multi-person facial expression recognition. *International Journal of Human-Computer Interaction*, 1–19. https://doi.org/10.1080/10447318.2024.2338614

Li, S., & Deng, W. (2022). Deep Facial Expression Recognition: A Survey. *IEEE Transactions on Affective*

Computing, *13*(3), 1195–1215. https://doi.org/10.1109/TAFFC.2020.2981446

Liu, Y., Dai, W., Feng, C., Wang, W., Yin, G., Zeng, J., & Shan, S. (2022). MAFW: A large-scale, multi-modal, compound affective database for dynamic facial expression recognition in the wild. *Proceedings of the 30th ACM International Conference on Multimedia*, 24–32. https://doi.org/10.1145/3503161.3548190

Manelis-Baram, L., Barami, T., Ilan, M., Meiri, G., Menashe, I., Soskin, E., Sofer, C., & Dinstein, I. (2025). Comparing three algorithms of automated facial expression analysis in autistic children: Different sensitivities but consistent proportions. *Molecular Autism*, *16*(1). https://doi.org/10.1186/s13229-025-00685-x

McKeown, G., & Allison, B. (2024). Are facial expression technologies tools for emotion recognition or social interaction analysis, or both? *Proceedings of the 11th International Conference on Affective Computing and Intelligent Interaction, ACII 2023, International Conference on Affective Computing and Intelligent Interaction (ACII): Proceedings*.

Mehu, M., & Scherer, K. R. (2015). Emotion categories and dimensions in the facial communication of affect: An integrated approach. *Emotion*, *15*(6), 798–811. https://doi.org/10.1037/a0039416

Meng, Z., Liu, P., Cai, J., Han, S., & Tong, Y. (2017). Identity-Aware Convolutional Neural Network for Facial Expression Recognition. *2017 12th IEEE International Conference on Automatic Face & Gesture Recognition (FG 2017)*, 558–565. https://doi.org/10.1109/FG.2017.140

Mollahosseini, A., Hasani, B., & Mahoor, M. H. (2019). AffectNet: A database for facial expression, valence, and arousal computing in the wild. *IEEE Transactions on Affective Computing*, *10*(1), 18–31. https://doi.org/10.1109/TAFFC.2017.2740923

Nesselroade, J. R., McArdle, J. J., Aggen, S. H., & Meyers, J. M. (2013). Dynamic factor analysis models for representing process in multivariate time-series. In *Modeling intraindividual variability with repeated measures data* (pp. 235–265). Psychology Press.

Ollero, M. J. F., Estrada, E., Hunter, M. D., & Cáncer, P. F. (2023). Characterizing affect dynamics with a damped linear oscillator model: Theoretical considerations and recommendations for individual-level applications. *Psychological Methods*. https://doi.org/10.1037/met0000615

Plutchik, R. (2001). The nature of emotions. *American Scientist*, *89*(4), 344–350. https://doi.org/10.1511/2001.4.344

Pons, P., & Latapy, M. (2005). Computing communities in large networks using random walks. In Pi. Yolum, T. Güngör, F. Gürgen, & C. Özturan (Eds.), *Computer and information sciences - ISCIS 2005* (pp. 284–293). Springer Berlin Heidelberg. https://doi.org/10.1007/11569596_31

Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021). Learning transferable visual models from natural language

supervision. *Proceedings of the 38th International Conference on Machine Learning*, 8748–8763.

Russell, J. A. (1980). A circumplex model of affect. *Journal of Personality and Social Psychology*, *39*(6), 1161–1178. https://doi.org/10.1037/h0077714

Russell, J. A. (2003). Core affect and the psychological construction of emotion. *Psychological Review*, *110*(1), 145–172. https://doi.org/10.1037/0033-295X.110.1.145

Russell, J. A., & Mehrabian, A. (1977). Evidence for a three-factor theory of emotions. *Journal of Research in Personality*, *11*(3), 273–294. https://doi.org/10.1016/0092-6566(77)90037-X

Stewart, P. A., Waller, B. M., & Schubert, J. N. (2009). Presidential speechmaking style: Emotional response to micro-expressions of facial affect. *Motivation and Emotion*, *33*(2), 125–135. https://doi.org/10.1007/s11031-009-9129-1

Tomašević, A., Golino, H., & Christensen, A. (2026). transforEmotion: An open-source R package for emotion analysis using transformer-based generative AI models. *Computational Communication Research*, *8*(2), 1. https://doi.org/10.5117/CCR2026.2.2.TOMA

Tong, E. M. W., Bishop, G. D., Enkelmann, H. C., Why, Y. P., Diong, S. M., Khader, M., & Ang, J. (2005). The Use of Ecological Momentary Assessment to Test Appraisal Theories of Emotion. *Emotion*, *5*(4), 508–512. https://doi.org/10.1037/1528-3542.5.4.508

Traag, V. A., Waltman, L., & Eck, N. J. van. (2019). From louvain to leiden: Guaranteeing well-connected communities. *Scientific Reports*, *9*(1), 5233. https://doi.org/10.1038/s41598-019-41695-z

Wang, W., Sebe, N., & Lepri, B. (2022). *Rethinking the Learning Paradigm for Facial Expression Recognition* (arXiv:2209.15402). arXiv. https://arxiv.org/abs/2209.15402

Appendix A: Population Perturbation View of LOMM Recovery

This appendix formalizes the population argument behind the Discussion. The result is not a theorem about the fitted DynEGA algorithm, its community-detection implementation, or the exact behavior of GraphicalVAR or GIMME software. It is a theorem about the simplified LOMM population object that any dynamic dimensional-recovery method receives: observed score channels in the same latent block share coordinated changes because they are noisy indicators of the same activity-scaled oscillator.

A.1 Simplified Population LOMM

Consider a population LOMM with m latent dimensions, k observed indicators per dimension, and $p = mk$ observed score channels. Let $B_1, \dots, B_m$ denote the true indicator blocks. For indicator $j \in B_f$, assume the equal-loading, equal-error measurement model

$$Y_j(t) = \lambda Z_f(t) + \varepsilon_j(t),$$

where

$$Z_f(t) = a_f(t) X_f(t)$$

is the activity-scaled latent oscillator for dimension f, $\lambda > 0$, and $\varepsilon_j(t)$ has mean zero and variance $\psi = s_\varepsilon^2$. Measurement errors are independent across indicators and time, and the centered latent processes Z_f are mutually independent with finite variance. The implemented simulation allows loading jitter and item-level error jitter, but this symmetric case isolates the block geometry that defines the known-truth recovery target.

A.2 Dynamic Coordination Graph

Define the lag-h score change as

$$\Delta_h Y_j(t) = Y_j(t+h) - Y_j(t),$$

and define a population dynamic coordination graph by

$$W^{\Delta}_{j\ell} = \left|\mathrm{Cor}(\Delta_h Y_j, \Delta_h Y_\ell)\right|.$$

For two distinct indicators $j, \ell \in B_f$,

$$\Delta_h Y_j(t) = \lambda \Delta_h Z_f(t) + \Delta_h \varepsilon_j(t),$$

so

$$\mathrm{Cov}(\Delta_h Y_j, \Delta_h Y_\ell) = \lambda^2 \,\mathrm{Var}(\Delta_h Z_f)$$

and

$$\mathrm{Var}(\Delta_h Y_j) = \lambda^2 \,\mathrm{Var}(\Delta_h Z_f) + 2\psi.$$

Therefore the within-block dynamic edge is

$$w_f^{\Delta} = \frac{\theta_f^{\Delta}}{1 + \theta_f^{\Delta}}, \qquad \theta_f^{\Delta} = \frac{\lambda^2 \,\mathrm{Var}(\Delta_h Z_f)}{2\psi}.$$

For $j \in B_f$ and $\ell \in B_g$ with $f \neq g$, independence of the latent processes and measurement errors gives zero cross-block covariance. Thus the population dynamic coordination graph has block-clique form:

$$W^{\Delta} = \mathrm{blockdiag}\left(w_1^{\Delta}(J_k - I_k), \ldots, w_m^{\Delta}(J_k - I_k)\right).$$

This is the population object most directly aligned with dynamic dimensional recovery: variables in the same block share the same latent change signal, while variables in different blocks do not.

A.3 Partial-Correlation Graph

The GraphicalVAR comparator is converted into dimensional partitions by applying community detection to a contemporaneous partial-correlation matrix. To see how this representation changes the same block structure, consider the population covariance matrix inside block B_f. Let

$$S_f = \mathrm{Var}(Z_f), \qquad \theta_f = \frac{\lambda^2 S_f}{\psi}.$$

The block covariance matrix is

$$\Sigma_f = \lambda^2 S_f \mathbf{1}_k \mathbf{1}_k^\top + \psi I_k = \psi(I_k + \theta_f \mathbf{1}_k \mathbf{1}_k^\top).$$

Using the Sherman-Morrison identity,

$$\Sigma_f^{-1} = \psi^{-1}\left[I_k - \frac{\theta_f}{1+k\theta_f}\mathbf{1}_k\mathbf{1}_k^\top\right].$$

The off-diagonal precision entry is

$$(\Sigma_f^{-1})_{j\ell} = -\psi^{-1}\frac{\theta_f}{1+k\theta_f}, \qquad j \neq \ell,$$

and the diagonal precision entry is

$$(\Sigma_f^{-1})_{jj} = \psi^{-1}\frac{1+(k-1)\theta_f}{1+k\theta_f}.$$

The absolute partial correlation between two indicators in the same block is therefore

$$w_f^P = \frac{\theta_f}{1+(k-1)\theta_f}.$$

Because the full covariance matrix is block diagonal under the simplified population model, cross-block partial correlations are zero. The GraphicalVAR-style population graph preserves the ideal block support,

$$W^P = \text{blockdiag}\left(w_1^P(J_k - I_k), \ldots, w_m^P(J_k - I_k)\right),$$

but compresses the within-block signal. The corresponding marginal common-source edge would be

$$w_f^M = \frac{\theta_f}{1+\theta_f}.$$

For the same signal-to-noise ratio θ,

$$\frac{w^P(\theta)}{w^M(\theta)} = \frac{1+\theta}{1+(k-1)\theta}.$$

When $k > 2$ and $\theta > 0$, $w^P(\theta) < w^M(\theta)$. As $\theta \to \infty$, $w^P(\theta) \to 1/(k-1)$, whereas $w^M(\theta) \to 1$. Thus partial correlation preserves the ideal block support but attenuates the shared-source signal that defines a reflective measurement block.

A.4 Perturbation Margin

The preceding graphs are population graphs. Estimated graphs can be written as

$$\widehat{W} = W + E,$$

where E includes sampling error, derivative-estimation error, model-estimation error, finite-T and finite-N error, regularization, convergence behavior, and other implementation-specific perturbations. For a block-clique population graph

$$W = \text{blockdiag}\left(w_1(J_k - I_k), \ldots, w_m(J_k - I_k)\right),$$

let $w_{\min} = \min_f w_f$. Suppose every off-diagonal edge is estimated within δ of its population value:

$$\max_{j \neq \ell} |\widehat{W}_{j\ell} - W_{j\ell}| \leq \delta.$$

For any node $j \in B_f$, the estimated total connection to its true block is at least

$$(k-1)(w_f - \delta),$$

whereas its estimated total connection to any wrong block is at most

$$k\delta.$$

Thus each node has stronger total connection to its true block than to any wrong block whenever

$$(k-1)(w_{\min} - \delta) > k\delta,$$

or equivalently

$$\delta < \frac{k-1}{2k-1} w_{\min}.$$

This is the elementary perturbation statement. Recovery is favored when the estimation perturbation is small relative to the weakest within-block signal. Because the partial-correlation edge $w^P(\theta)$ is smaller than the marginal shared-source edge $w^M(\theta)$ for $k > 2$, the same perturbation has less margin under a partial-correlation representation.

A.5 Link to Oscillator Parameters

The signal-to-noise ratios above become LOMM-specific when $Z_f(t) = a_f(t)X_f(t)$ is tied to the latent oscillator. In the stable continuous-time case with constant activity a_f,

$$\ddot{X}_f(t) = \eta_f X_f(t) + \zeta_f \dot{X}_f(t) + q_f(t), \qquad \eta_f < 0,\ \zeta_f < 0.$$

Under the stationary linear approximation,

$$\mathrm{Var}(X_f) = \frac{\sigma_{qf}^2}{2\eta_f \zeta_f},$$

which is positive because both η_f and ζ_f are negative. Therefore

$$\mathrm{Var}(Z_f) = a_f^2 \,\mathrm{Var}(X_f) = a_f^2 \frac{\sigma_{qf}^2}{2\eta_f \zeta_f}.$$

The contemporaneous block signal is controlled by

$$\theta_f = \frac{\lambda^2 a_f^2 \sigma_{qf}^2}{2\eta_f \zeta_f \psi}.$$

For the dynamic coordination graph,

$$\mathrm{Var}(\Delta_h Z_f) = 2a_f^2 \,\mathrm{Var}(X_f)\{1 - \rho_f(h)\},$$

SO

$$\theta_f^{\Delta}(h) = \frac{\lambda^2 a_f^2 \operatorname{Var}(X_f)\{1 - \rho_f(h)\}}{\psi}.$$

In the underdamped case,

$$\rho_f(h) = e^{\zeta_f h/2}\left[\cos(\omega_f h) - \frac{\zeta_f}{2\omega_f}\sin(\omega_f h)\right], \qquad \omega_f = \sqrt{-\eta_f - \frac{\zeta_f^2}{4}}.$$

Thus the perturbation margin depends on the same quantities varied in the simulation design: loading strength λ, measurement-error variance $\psi = s_\varepsilon^2$, activity a_f, shock variance σ_{qf}^2, position feedback η_f, velocity feedback ζ_f, lag h, and number of indicators k.

A.6 Interpretation

This derivation explains the target-alignment claim in the Discussion. In the simplified population LOMM, the true dimensions are disconnected weighted cliques or, equivalently, rank-one dynamic subspaces. Finite samples and estimator behavior perturb that graph. Recovery is favored when the within-block signal remains large relative to the perturbation.

DynEGA is aligned with this target because it is used here to recover groups of variables whose changes are coordinated over time. GraphicalVAR remains a reasonable network comparator, but the contemporaneous partial-correlation representation conditions on other indicators in the same reflective block and therefore compresses the shared-source signal. GIMME is not part of this theorem because it addresses a broader idiographic time-series problem with group-level, subgroup-level, and individual-level paths; in this study it is best interpreted as a flexible but computationally bounded feasibility comparator.

Appendix B: Empirical Realism Numeric Summary

Table 3: Numeric summary for the empirical realism check. The main Results section reports these calls visually in Figure 4; this appendix table gives the corresponding medians and discrepancy components.

Readout	Metric	Empirical median	LOMM median	Scaled median diff.	Scaled Wasserstein	Max discrepancy	Status
Channels	Channel IQR	0.020	0.064	0.76	1.06	1.06	Fail
Channels	Derivative IQR	0.014	0.013	0.01	0.23	0.23	Pass
Channels	Lag-1 ACF	0.521	0.624	0.33	0.28	0.33	Partial
Channels	Run-length entropy	0.785	0.746	0.30	0.29	0.30	Pass
Channels	Spectral entropy	0.755	0.685	0.42	0.50	0.50	Partial
Channels	Effective rank	5.114	5.959	0.63	0.54	0.63	Partial
Community scores	Channel IQR	0.247	0.832	1.40	1.48	1.48	Fail
Community scores	Derivative IQR	0.159	0.294	0.52	0.50	0.52	Partial
Community scores	Lag-1 ACF	0.540	0.745	0.75	0.56	0.75	Partial
Community scores	Run-length entropy	0.785	0.756	0.22	0.22	0.22	Pass
Community scores	Spectral entropy	0.751	0.637	0.69	0.69	0.69	Partial
Community scores	Effective rank	2.371	2.803	0.76	0.77	0.77	Partial